



Hardware acceleration for PyTorch – a demonstrator

with Embecosm


The team

MEng students at the University of Southampton.

- James Betson (electronic engineering)
- Lorinc Boer (computer science)
- Reuben Leanage (electronic engineering)
- Tao Zeng (electronic engineering)
- Ze Yong (electronic engineering)
- AoYang Leng (electronic engineering)

PyTorch

- Python API for building **neural networks**.
- Built on top of C++ using pybind.
- Used by Meta, Apple, TikTok, Microsoft...
- **Quick** and **easy**.



A fly in the ointment?

- PyTorch has great support for CPUs and GPUs.



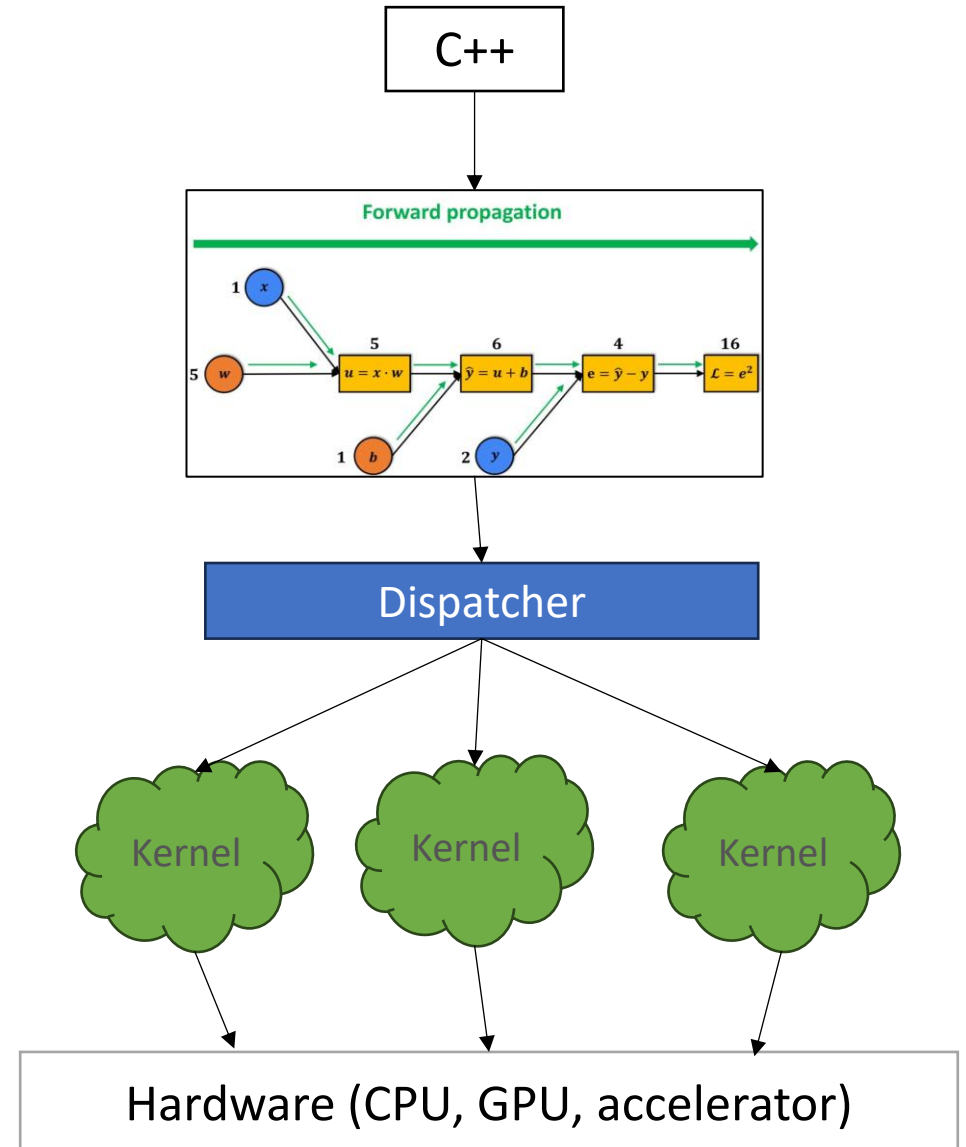
Parallel SIMD
compute

Fast matrix
computation

- Support for custom hardware? What about a bespoke ASIC or FPGA design?

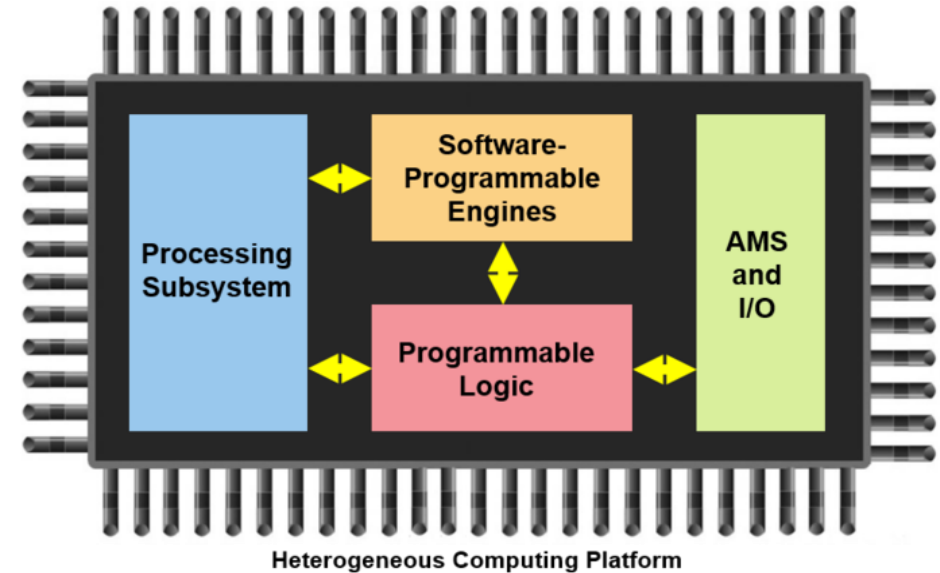
A fly in the ointment?

- Hardware dependent **kernels**



Custom accelerators

- Improved Performance.
- Specialized hardware.
- Unlimited flexibility.

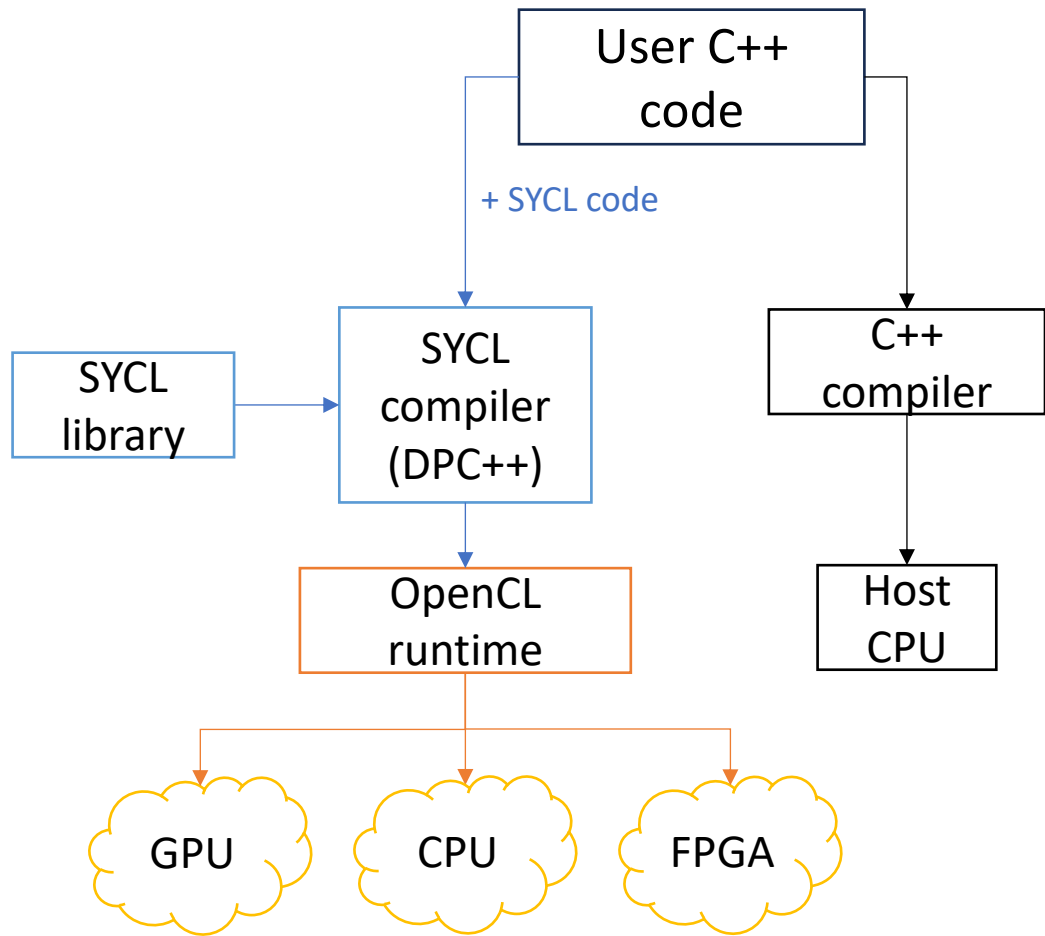


The project


“One of Pytorch’s current weaknesses is how **difficult it is to add support for new accelerators** in Eager mode. As it stands Pytorch defines several thousand core operations, each of which must be implemented individually for any new hardware – **an extremely large undertaking.**”

- Investigate how this can be done.
- Build a proof-of-concept demonstrator.

SYCL: How does it work?




- C++ programming model
- Generic hardware-agnostic (abstraction)




- Heterogeneous compute framework
- Interfaces with hardware

Dispatch Comparison



Typical PyTorch offload



(Theoretical) PyTorch SYCL offload

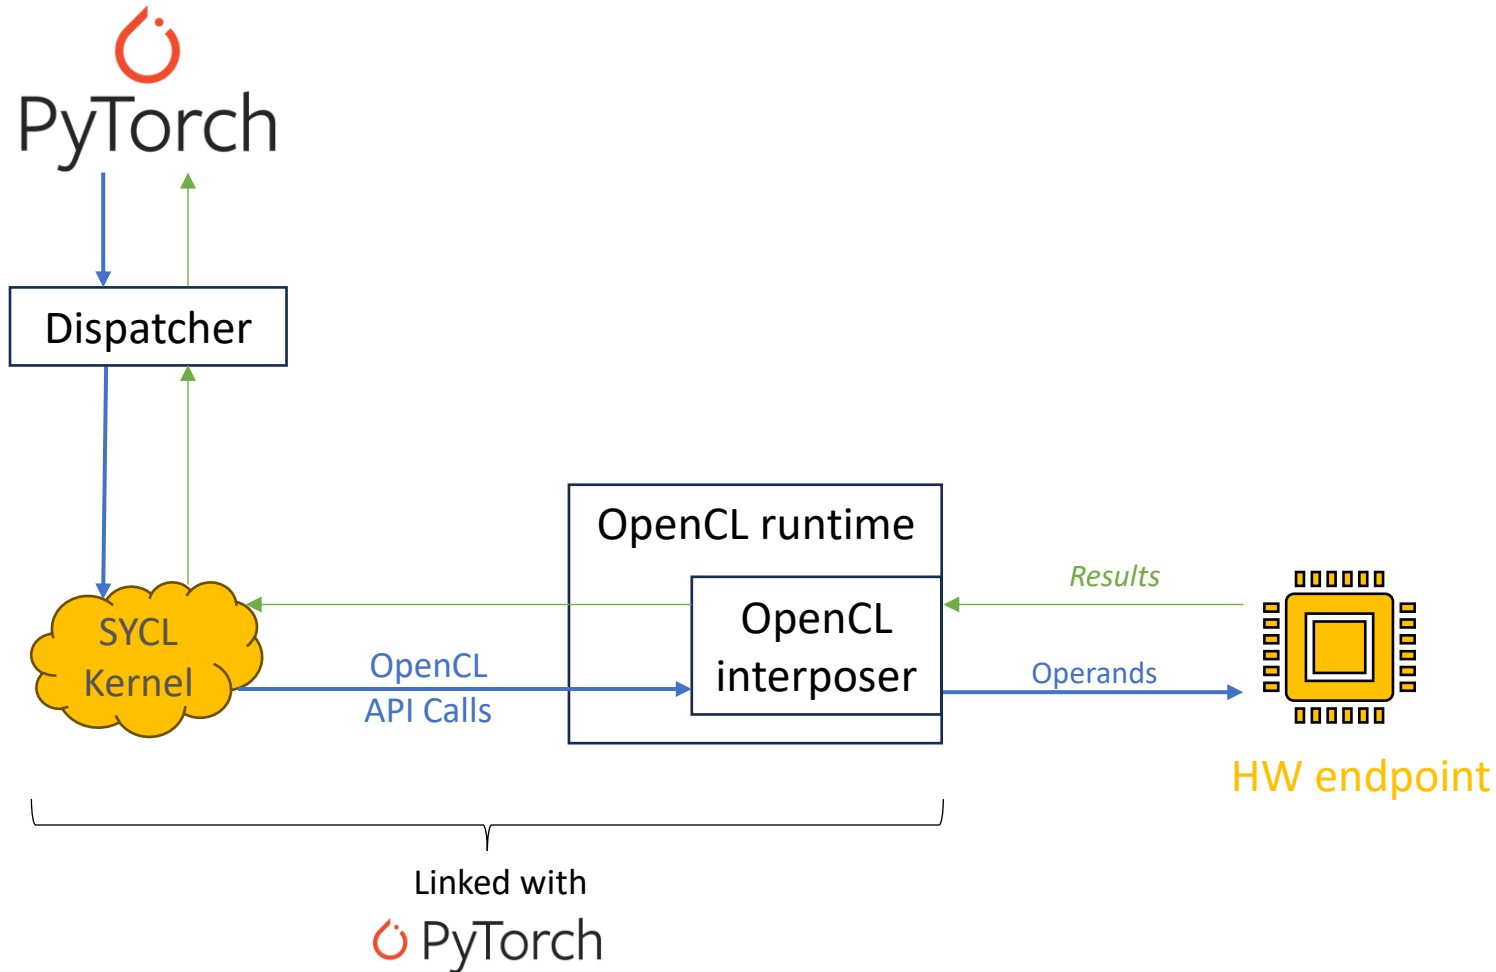
C++ extensions

- PyTorch tool (pybind).
 - Only supports a standard g++ compilation process.


Solution: write standard C++ kernels

- Call SYCL functions as 'extern'.
- Compile SYCL functions separately into a **shared object**.
- Link: `'export LD_PRELOAD=$(pwd)/build/libsycl_library.so'`

Our Implementation



- C++ programming model
- Generic hardware-agnostic (abstraction)



- Heterogeneous compute framework
- Interfaces with hardware

Interposer – why?

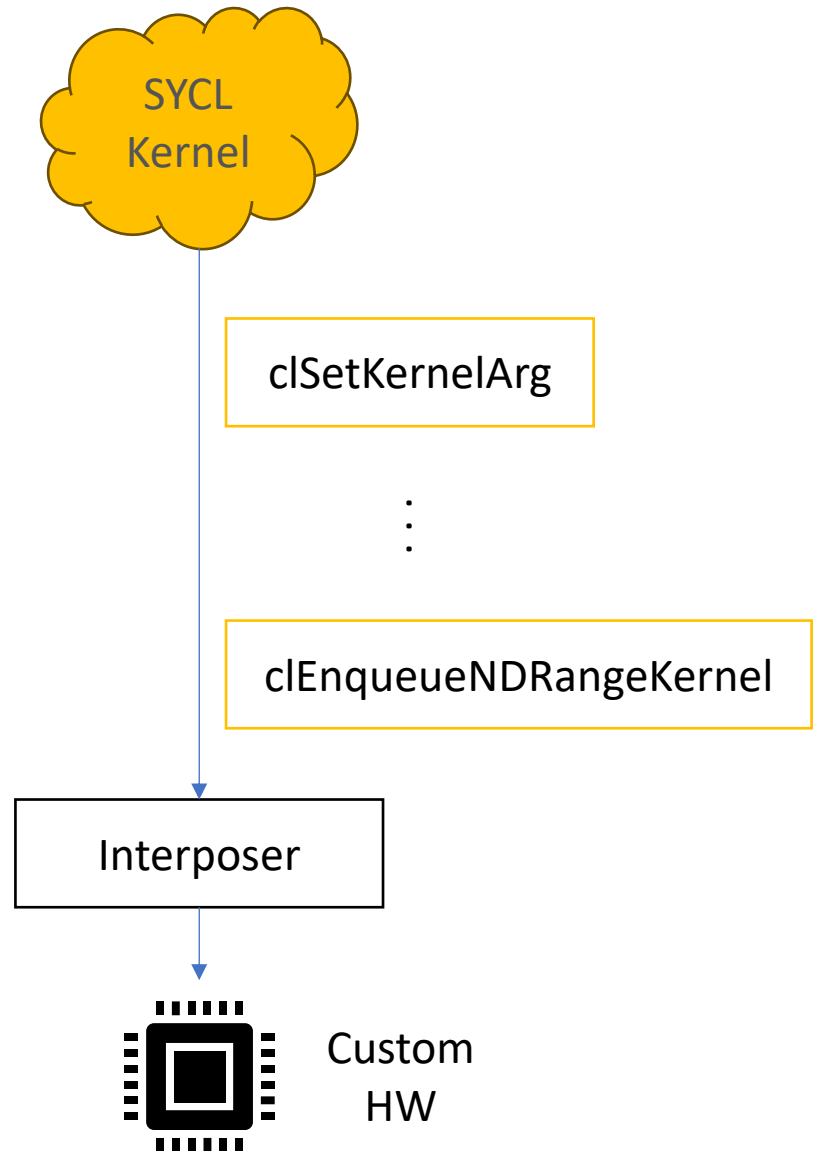
- oneAPI HAL: struggled to bring up. Worried about overhead.
- OpenCL device: would take too long. Lots of overhead.

Solution: manually intercept OpenCL calls to *'sycl::default_selector_v'*.

- Extract operands and operation type.

Interposer – how?

1. `clSetKernelArg` – captures SYCL data buffer
2. ...
3. ...
4. ...
5. ...
6. ...
7. ...
8. `clEnqueueNDRangeKernel` – trigger computation



Why is this better?

- No more re-writing PyTorch source code!
- User engineering effort has changed:

Before	After
Re-write potentially 1000s of kernels – hours of work	Use PyTorch as is – easy
Repeat this process and re-compile every time hardware changes – tedious	Interface with OpenCL API calls. Easy to intercept data for your hardware to run

Why is this better

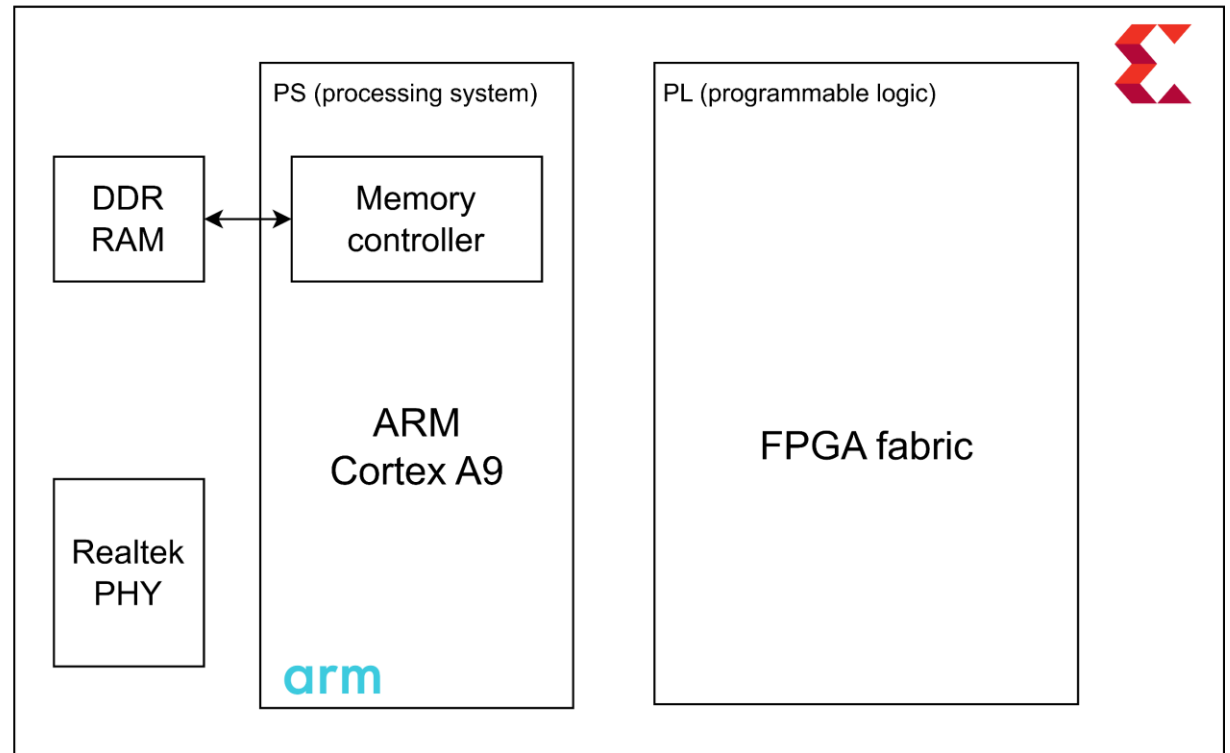
- Shifting user effort **away** from PyTorch.
- AI tools must keep up!

Our demonstrator

Linux x86 host



Xilinx Zynq-7000 SoC

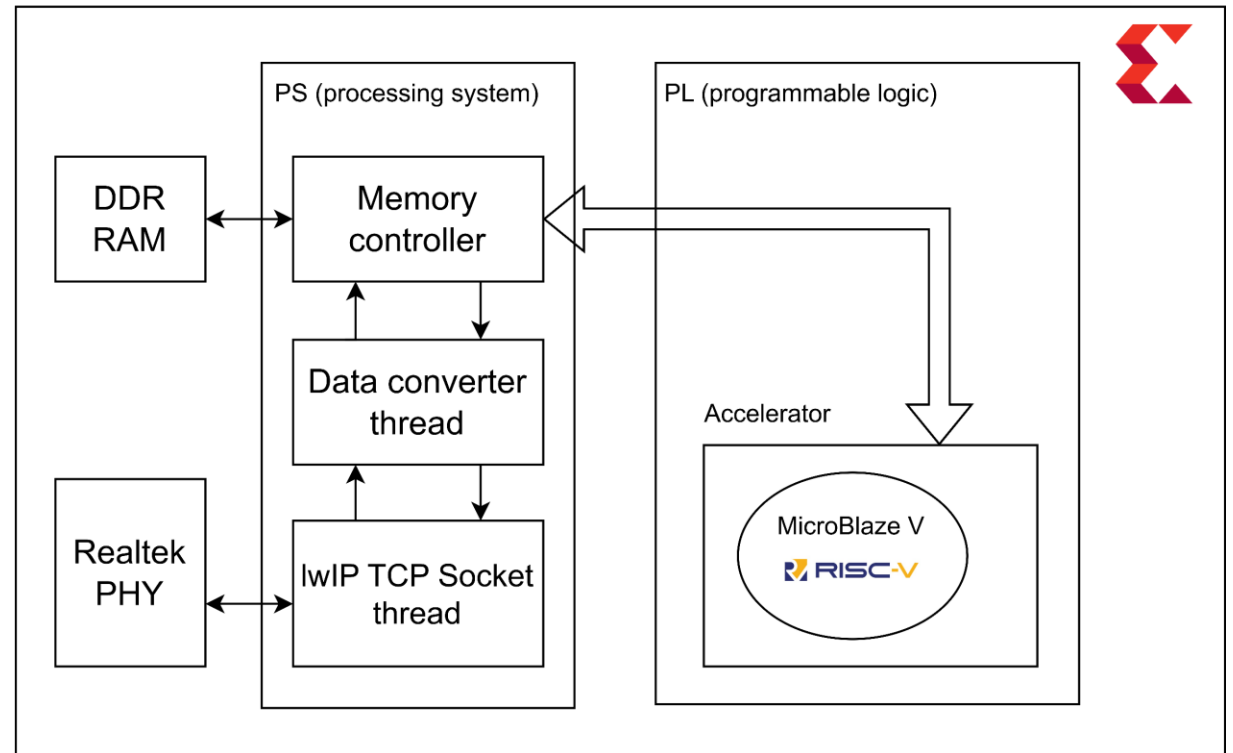


Our demonstrator

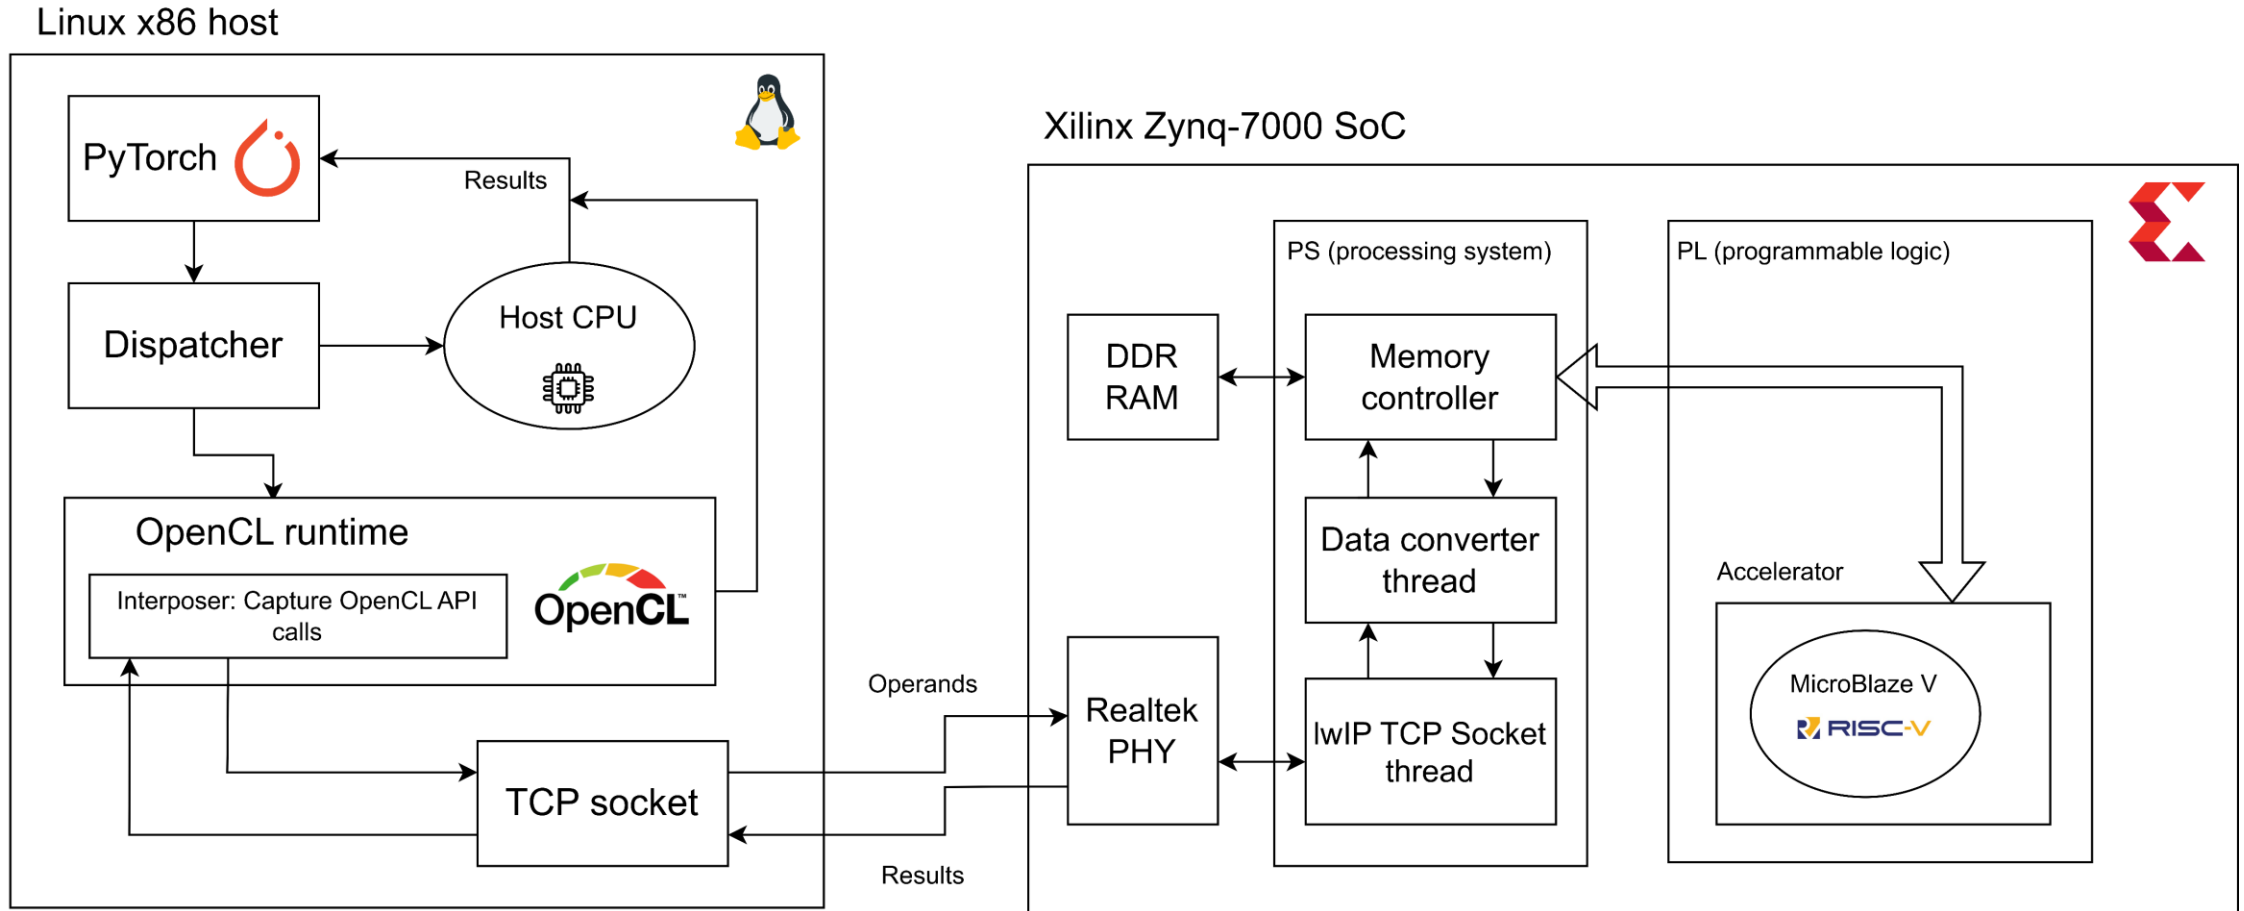
Linux x86 host



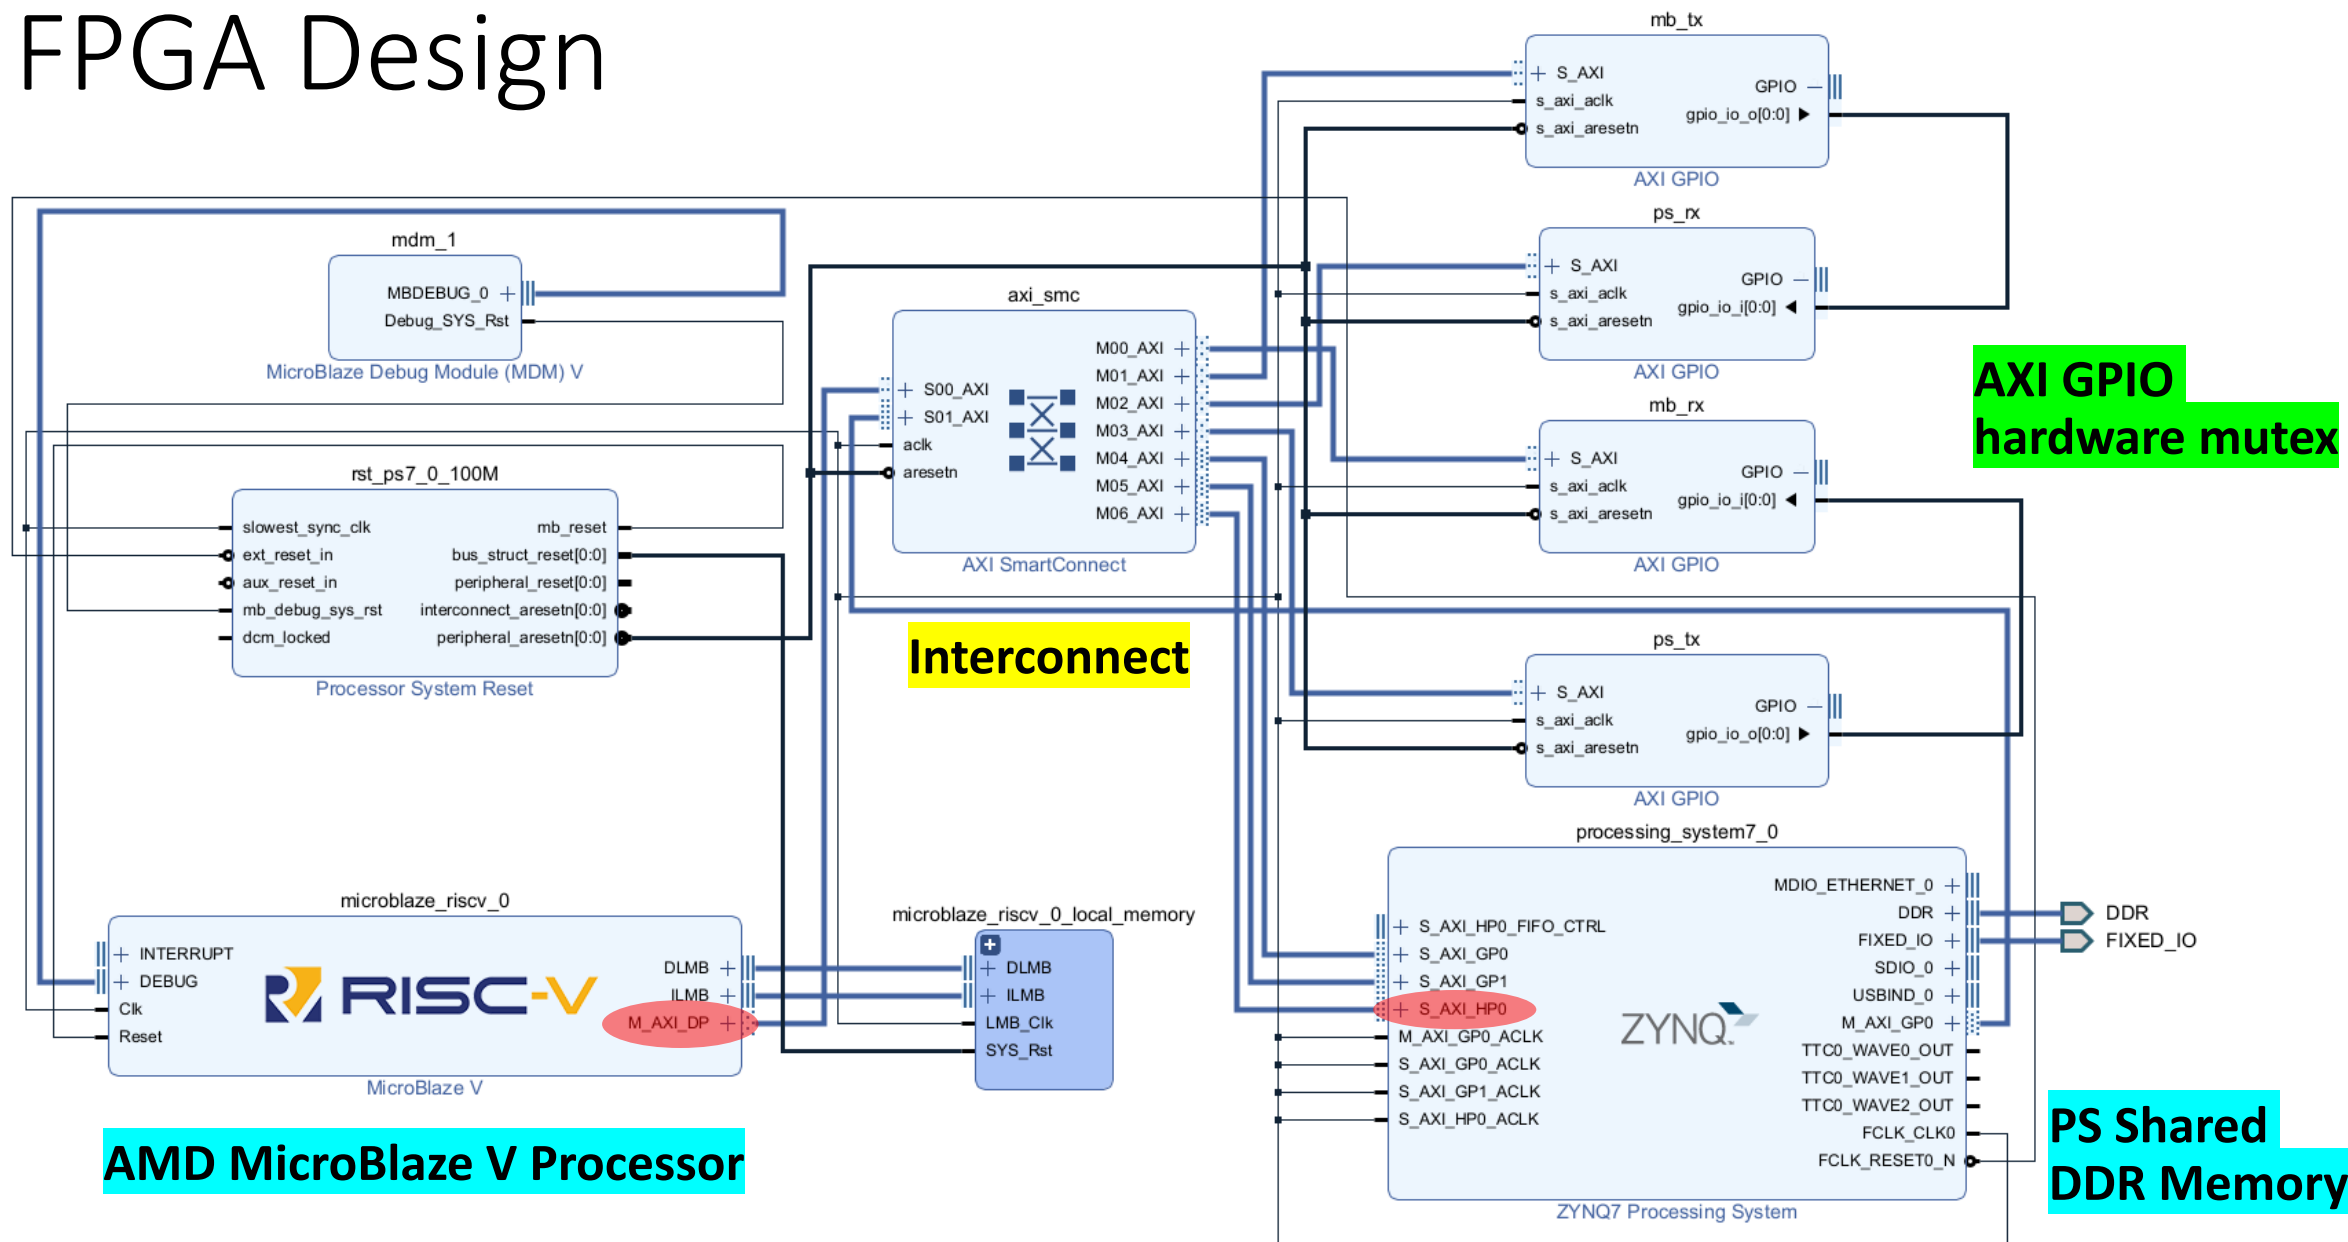
Xilinx Zynq-7000 SoC



Our demonstrator



FPGA Design



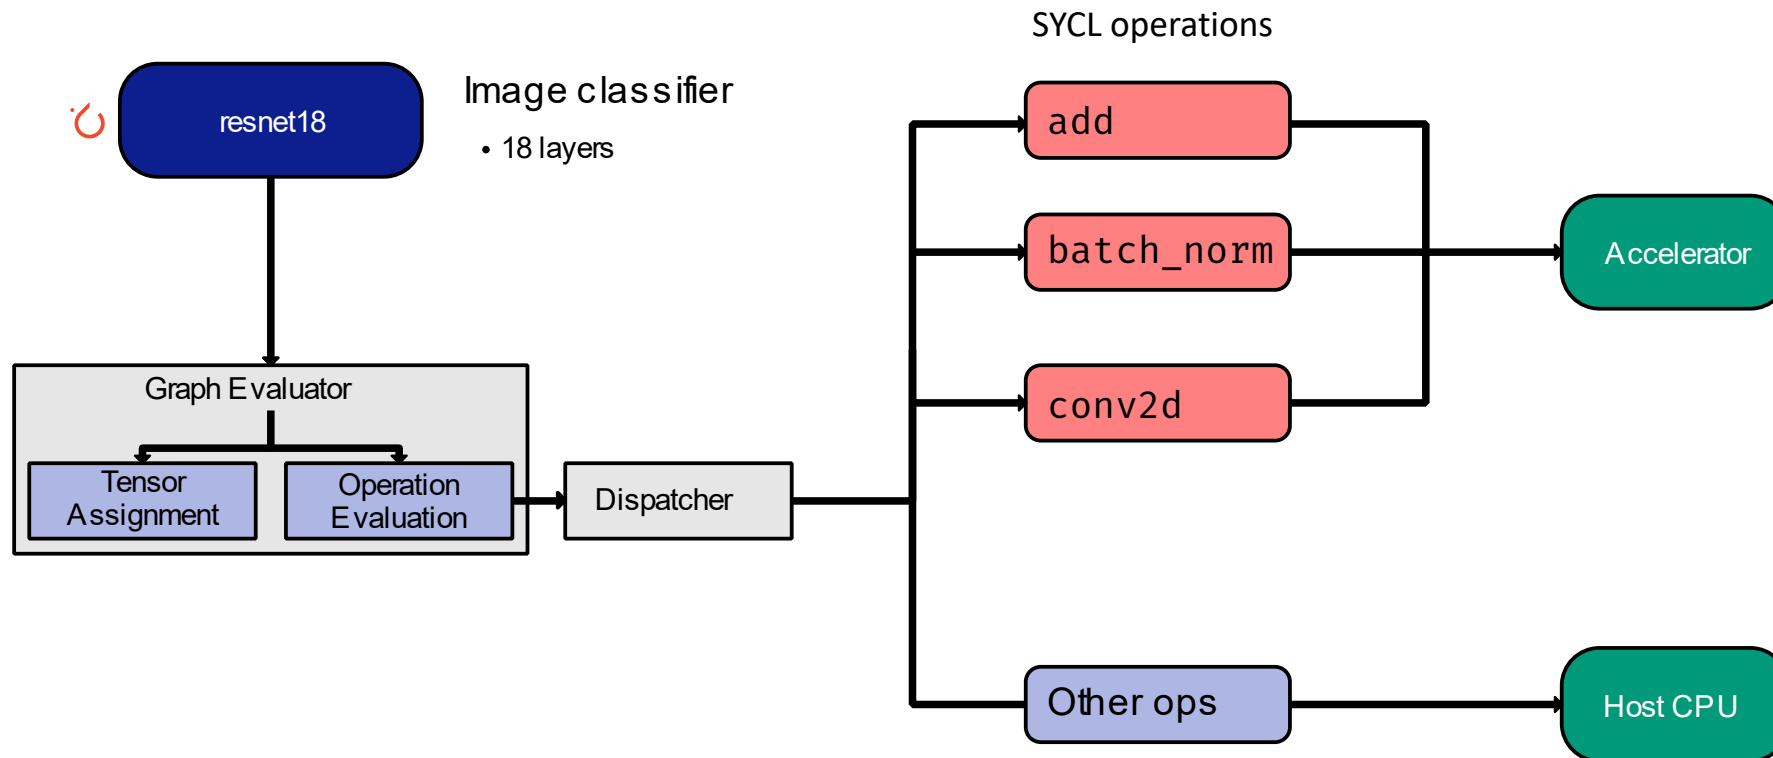
Proof-of-Concept

- RISC-V soft core = accelerator



- Show that we can offload compute from PyTorch to an accelerator
 - *Without* the user re-writing the PyTorch backends!
 - No re-writing for *new* accelerators – **scalable, hardware-agnostic**

Proof-of-concept



Results

```
# SYCL kernel computation
result_sycl = extension_cpp.ops.sycl_add_out(a, b, alpha, result_is_same_as_other)

# PyTorch computation
result_pytorch = torch.add(a, b, alpha=alpha)
```

```
(env) jbetson@James:~/SYCL_extensions$ python3 test_extension.py
[Python] a: tensor([[-0.0883,  0.3420, -1.6348],
                  [-0.6108,  0.6003, -0.8769],
                  [ 0.9649, -0.1926,  0.1148],
                  [-0.1952,  0.3025,  1.3122],
                  [-0.5610,  0.1900, -0.1660],
                  [ 0.9107, -0.3660,  0.0853]], dtype=torch.float64)
[Pytohn] b: tensor([[ 0.8843,  1.5818,  0.1156],
                  [ 0.9368,  0.7956, -0.0108],
                  [ 3.1580, -0.8736,  1.3865],
                  [ 1.4566,  0.3146, -1.7189],
                  [ 0.9152, -0.2367, -0.2758],
                  [ 1.5137, -0.2111,  1.8952]], dtype=torch.float64)
[Python] alpha: 2.0
```


```
ZE_LOADER_DEBUG_TRACE:Using Loader Library Path:
ZE_LOADER_DEBUG_TRACE:0 Drivers Discovered
[Interposer] clEnqueueNDRangeKernel called
[Interposer] Received 10 clSetKernelArg OpenCL API calls
[Interposer] 'add' detected
[Interposer] Operand A:
-0.0883101 0.342049 -1.63476 -0.610772 0.600292 -0.876878 0.964896
-0.165965 0.910662 -0.366011 0.0852656
[Interposer] Operand B:
0.884346 1.5818 0.115601 0.936776 0.79561 -0.0107679 3.15796 -0.873
1.51365 -0.211141 1.89521
[Interposer] ScaleA: 1
[Interposer] ScaleB: 2
[Interposer] Connected to client at 192.168.1.10:7
[Interposer] Sent 39 doubles (312 bytes) to client.
```

Results

```
----- In Python -----  
Result accelerated:  
tensor([[ 1.6804,  3.5057, -1.4036],  
        [ 1.2628,  2.1915, -0.8984],  
        [ 7.2808, -1.9398,  2.8879],  
        [ 2.7181,  0.9317, -2.1255],  
        [ 1.2693, -0.2833, -0.7176],  
        [ 3.9380, -0.7883,  3.8757]], dtype=torch.float64)  
Result PyTorch:  
tensor([[ 1.6804,  3.5057, -1.4036],  
        [ 1.2628,  2.1915, -0.8984],  
        [ 7.2808, -1.9398,  2.8879],  
        [ 2.7181,  0.9317, -2.1255],  
        [ 1.2693, -0.2833, -0.7176],  
        [ 3.9380, -0.7883,  3.8757]], dtype=torch.float64)  
PASSED: Results are identical
```

- For input tensors up to 512 doubles.
- Supports 'add' and 'batch_norm'.
- Insufficient time for 'conv2d' but SYCL support was built.

Case study: portability



Case study: performance

- Experimental results: ~5600 times slower for 'add' | ~9100 times slower for 'batch_norm'.

Table 7 – Final design performance results


Operation	Host CPU	MicroBlaze V (RISC-V)	MicroBlaze
Add	4.12 μ s	23,242.00 μ s	25,407.00 μ s
Batch_norm	10.87 μ s	99,213.64 μ s	102,160.08 μ s

Conclusion

What we did	What we didn't do
Hardware agnostic PyTorch offload	Accelerate PyTorch
Cheap FPGA hardware solution	Use oneAPI HAL to streamline hardware communication
Investigated OpenCL interposer methodology	

- Expansion to our suite of SYCL PyTorch kernels. Already work underway (SYCL C++ extensions).
- More performant hardware demonstrator. PCIe or CXL solution.
- Investigation into oneAPI HAL.

Video demo



[GitHub link](#)

Any questions?